

Ayush Saraf

📍 Vellore, India ✉ sarafa736@gmail.com ☎ +91 7004777571 🔗 [linkedin.com/in/saraf-ayush](https://www.linkedin.com/in/saraf-ayush) 📁 GitHub 🌐 Portfolio 📄 Coding Profile

Summary

Computer Science student with proficiency in Python, Java, C++, data structures and algorithms. Skilled in building practical projects involving NLP techniques, Machine Learning, LLMs, and AI Systems. Currently in pre-final year at VIT, focused on sharpening DSA skills for SDE roles.

Education

Vellore Institute of Technology <i>B. Tech in Computer Science (Core) CGPA: 8.88</i>	Vellore, India Aug 2023 – Present
Bhagwan Mahavir International School <i>Senior Secondary Percentage: 83.6%</i>	Gujarat, India 2023
L.P Savani School <i>Secondary Percentage: 88.86%</i>	Gujarat, India 2021

Skills

Languages: Python, Java, C/C++, JavaScript, TypeScript, SQL
Frameworks & Libraries: Scikit-learn, NumPy, Pandas, Matplotlib, Seaborn, OpenCV, Streamlit, LangChain, LangGraph
Platforms & Tools: Git, GitHub, VSCode, Cursor, Google Collab, Jupyter
Core Skills: OOP, Operating Systems, DSA
Soft Skills: Communication, Teamwork, Time Management, Self-directed Learning

Experience

AI Intern @ La Net Team Software Solution Pvt Ltd., Surat, Gujarat - Learned LangChain and LangGraph during the internship by building a RAG-based chatbot from scratch. - Built the retrieval pipeline (embeddings using Gemini API + vector search using Pinecone Database) over a repository of 50+ internal company documents, connecting it to Ollama for local LLM inference.	May 2026 – Jun 2026
AI/ML Intern @ SentiAid, Bengaluru, Karnataka - Cleaned and pre-processed a speech/text dataset (4,292 ISL sign videos, AI4Bharat) for an Indian Sign Language (ISL) translator using Python, key point detection, and pose-based animation. - Applied Pandas, Scikit-learn, and Regex to remove noise and align glosses with ASR outputs, improving training data quality.	May 2025 – Jul 2025

Projects

LangChain RAG Document Q&A System <i>LangChain, FAISS, Gemini, Groq, Streamlit, Python</i> - Built an end-to-end RAG pipeline integrating 4 AI APIs (Google Gemini Vision, Gemini Embeddings, Groq LLM, FAISS) to enable private-document Q&A. - Sped up embedding generation time by almost 10x using a parallel ThreadPoolExecutor (8 workers) and added incremental caching based on PDF timestamps so unchanged files re-index within 5 sec.	May 2026
Obsidian Linker <i>LangGraph, FAISS, Python, Pydantic, asyncio, Groq, Gemini</i> - Made a LangGraph agent with a 7-node pipeline that reads through an Obsidian vault and auto-generates <code>[[wiki-links]]</code> and <code>#tags</code> using Groq and Gemini as swappable LLM backends. - Set up concurrent LLM calls with <code>asyncio.gather</code>, and added MD5 hash checks so the pipeline skip notes that haven't changed, zero LLM calls on a re-run if nothing's new.	May 2026
Resume Feedback System <i>NLP, NLTK, TF-IDF, Python</i> - Built an NLP-powered ATS scoring system combining dynamic skill-matching density (~70% weight) and section/layout completeness (~30% weight) for 1–100% explainability. - Used NLTK sentence tokenizer to segment resumes into 5 categories, with a 2-layer text extraction pipeline. - Implemented TF-IDF vectorisation (up to 1000 features) with cosine similarity to evaluate keyword relevance and ATS-friendliness.	Oct 2025
Student Performance Predictor <i>ML, RandomForest, LLaMA, Python</i> - Trained a RandomForestClassifier on quiz/midterm/assignment data to predict pass/fail, got 94.5% accuracy. - Used Meta's LLaMA to turn the predictions into personalized feedback for each student.	Jul 2025

Certifications

Oracle Certified Foundations Associate, Java (1Z0-811) Oracle University	July 2026
---	-----------